

Dimensionality Reduction

Machine Learning and Deep Learning
Lesson #2

Multimedia DBs

- Many multimedia applications require efficient indexing in high-dimensions (time-series, images and videos, etc)
- Answering similarity queries in high-dimensions is a difficult problem due to “curse of dimensionality”
 - A solution is to use **Dimensionality reduction**
- The main idea: reduce the dimensionality of the space.
- Project the d-dimensional points in a k-dimensional space so that:
 - $k \ll d$
 - distances are preserved as well as possible
- Solve the problem in low dimensions

Multi-Dimensional Scaling (MDS)

- Map the items in a k-dimensional space trying to minimize the **stress**

$$stress = \sqrt{\frac{\sum_{i,j} (\hat{d}_{i,j} - d_{i,j})^2}{\sum_{i,j} d_{i,j}^2}} \text{ with}$$

$$d_{i,j} = |o_j - o_i|$$

$$\hat{d}_{i,j} = |\hat{o}_j - \hat{o}_i|$$

- **Steepest Descent algorithm:**
 - Start with an assignment
 - Minimize stress by moving points
- But the running time is $O(N^2)$ and $O(N)$ to add a new item

Embeddings

- Given a metric distance matrix D , embed the objects in a k -dimensional vector space using a mapping F such that
 - $D(i,j)$ is close to $D'(F(i),F(j))$
- Two types of mapping according to distances (Embedding):
 - **Isometric mapping:**
 - $D'(F(i),F(j)) = D(i,j)$
 - **Contractive mapping:**
 - $D'(F(i),F(j)) \leq D(i,j)$

Where d' is some L_p measure

- **Two types of embeddings according to warping technique**
 - *Linear* -> data points are projected by a linear transformation (PCA)
 - *Non linear* -> data points are projected non linearly (Laplacian ISOMAP, T-sne)



PCA Algorithm

- PCA algorithm:
 - 1. $X \leftarrow$ Create $N \times d$ data matrix, with one row vector x_n per data point
 - 2. X subtract mean \bar{x} from each row vector x_n in X
 - 3. $\Sigma \leftarrow$ covariance matrix of X
 - Find eigenvectors and eigenvalues of Σ
 - PC's \leftarrow the M eigenvectors with largest eigenvalues

Geometric Rationale of PCA

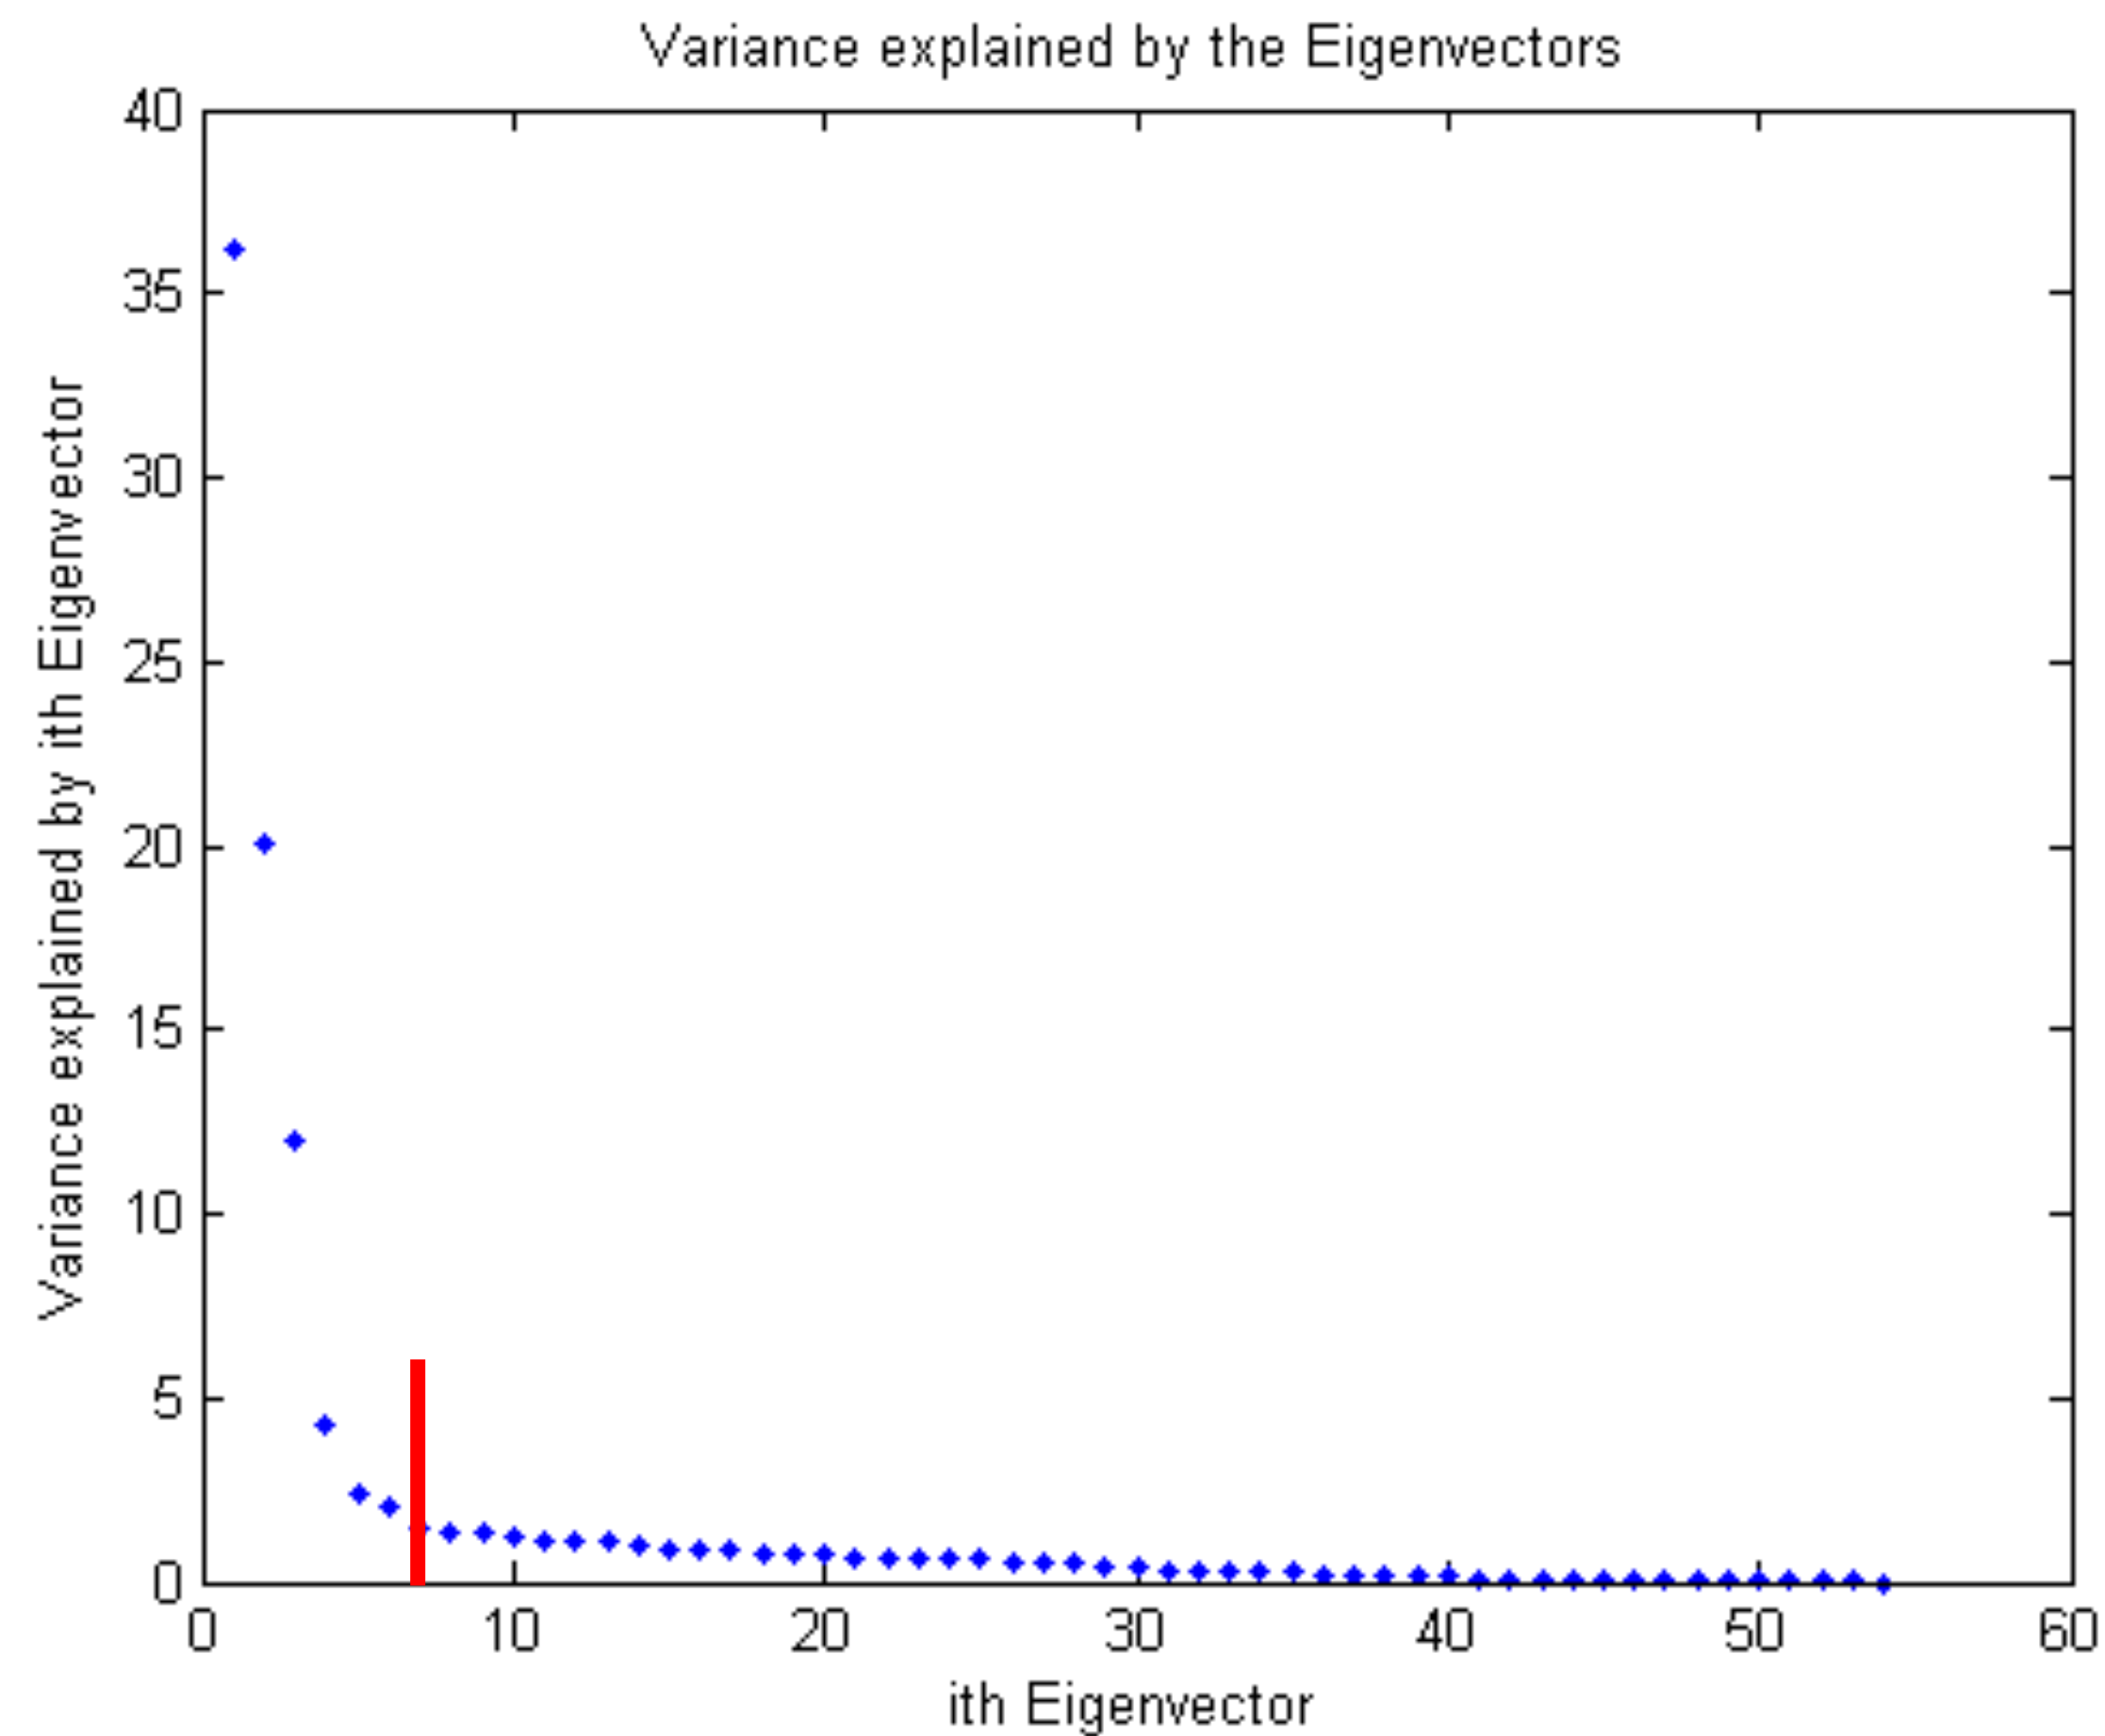
- objective of PCA is to rigidly rotate the axes of this p-dimensional space to new positions (principal axes) that have the following properties:
- ordered such that principal axis 1 has the highest variance, axis 2 has the next highest variance, , and axis p has the lowest variance
- covariance among each pair of the principal axes is zero (the principal axes are uncorrelated).

PCA principal AXIS

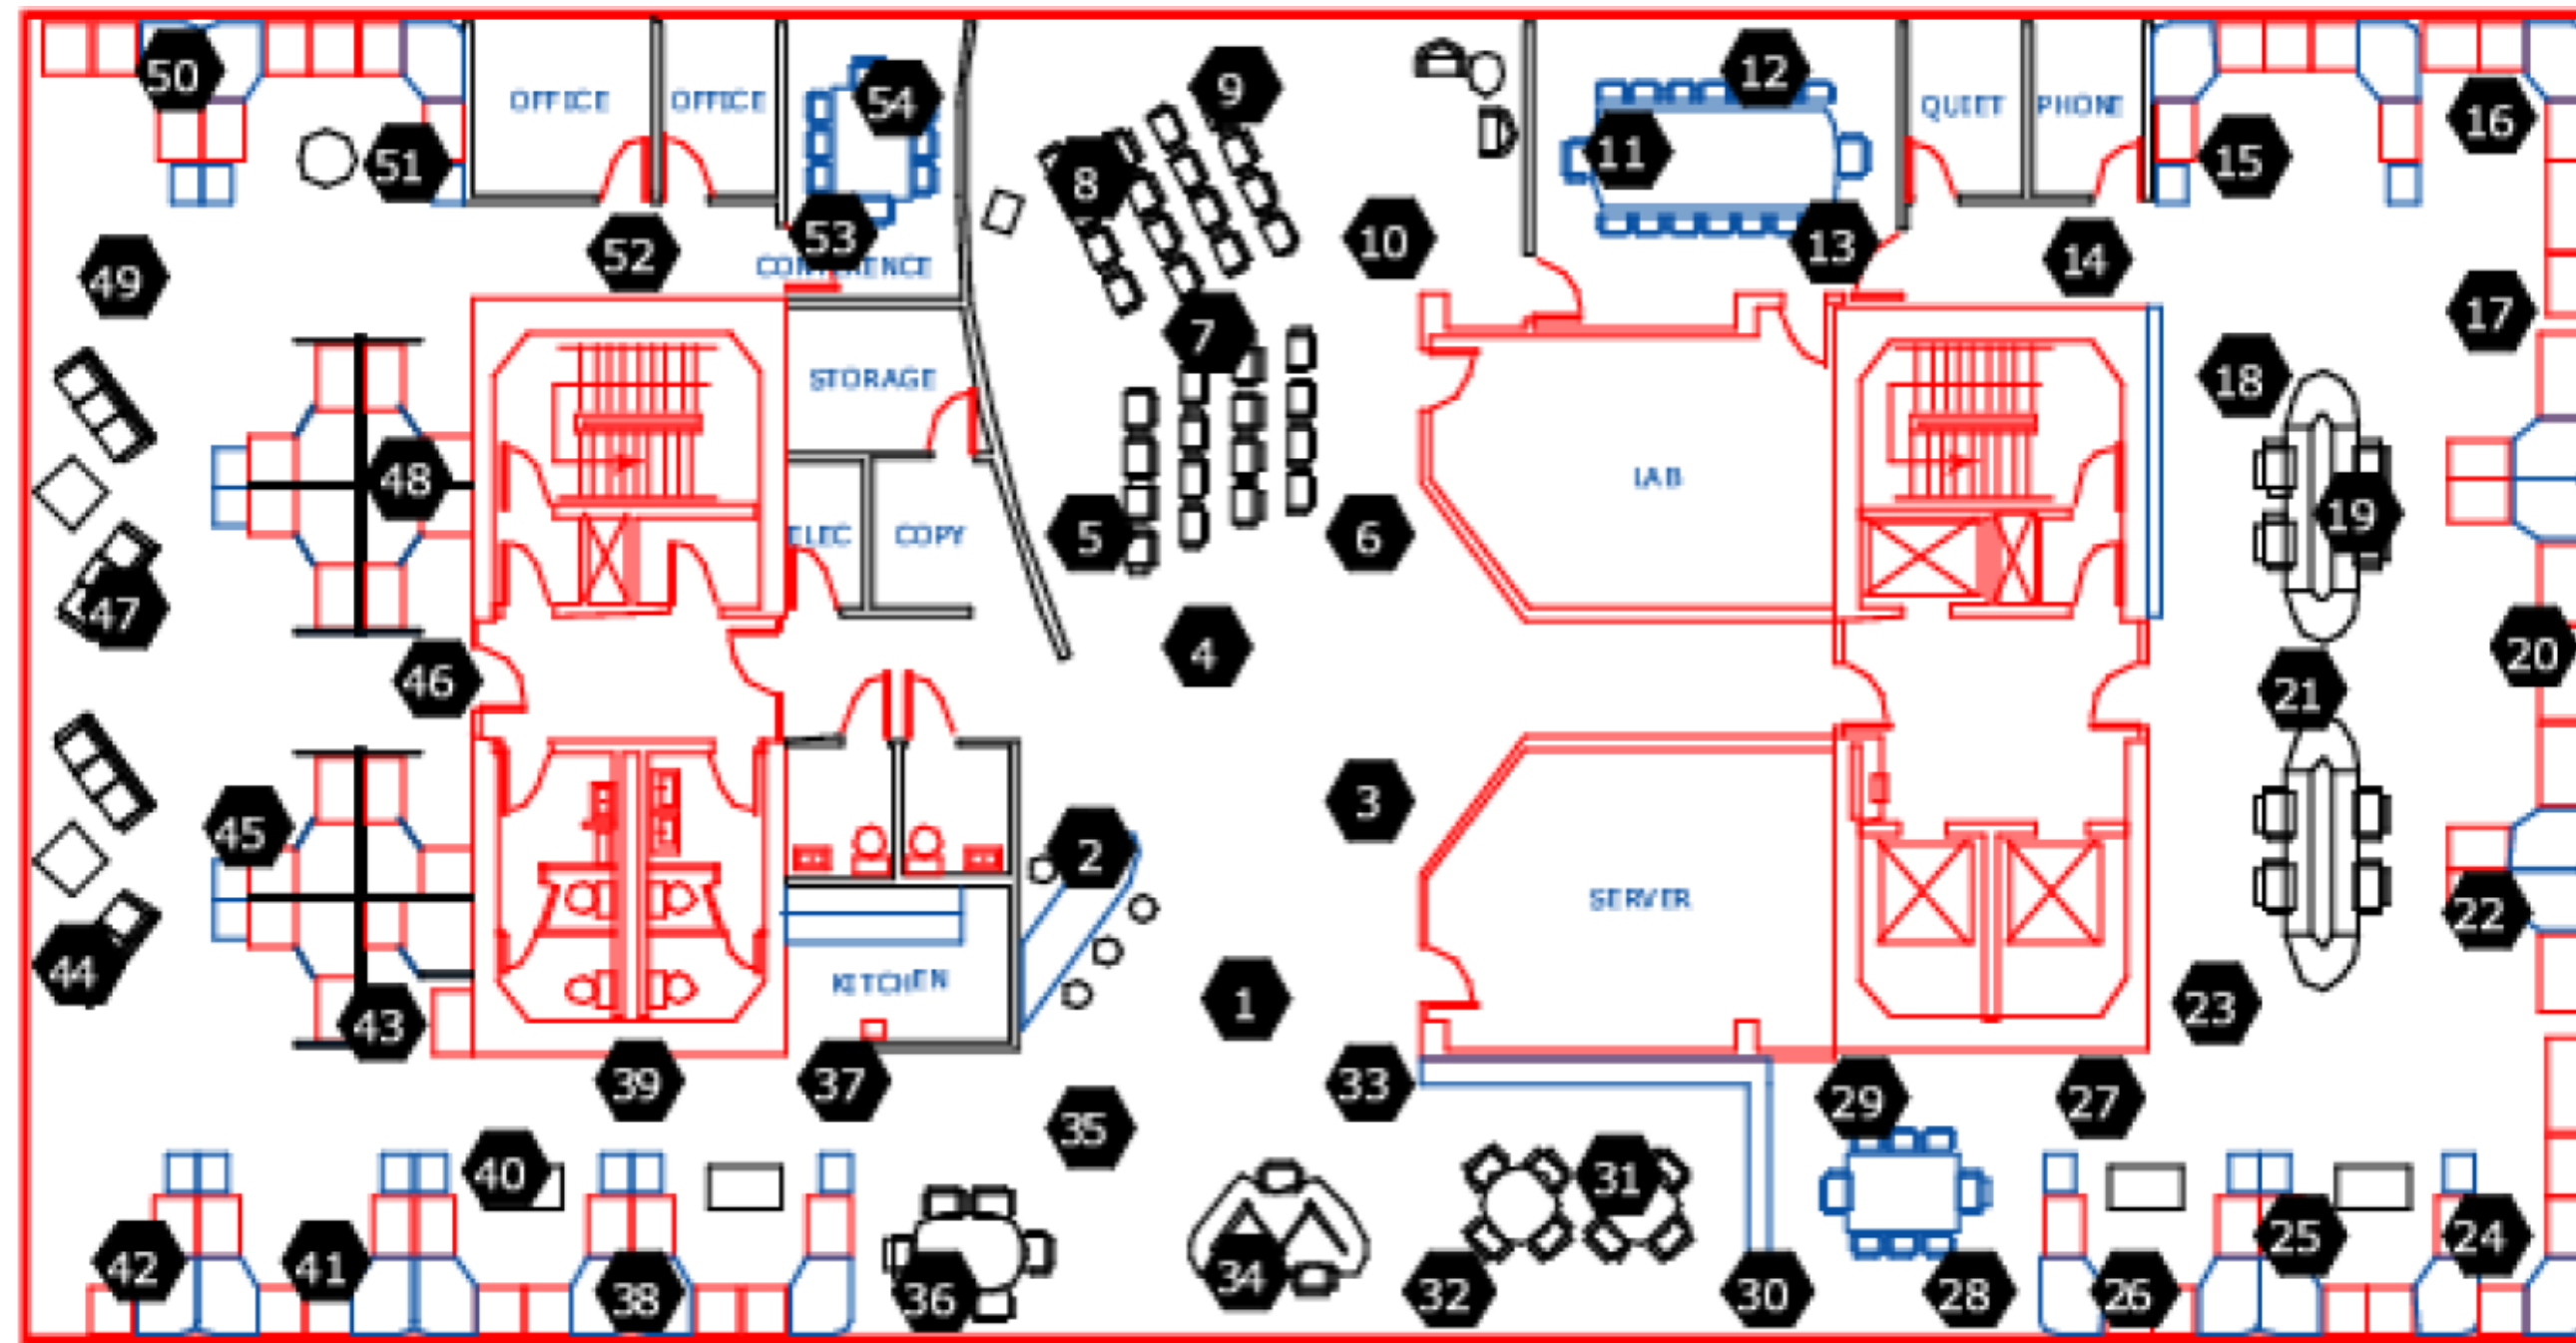


How many components?

- Check the distribution of eigen-values
- Take enough eigenvectors to cover 80-90% of the variance




Example Sensor networks




Sensors in Intel Berkeley Lab

Pairwise link quality vs. distance



PCA in action

- Given a 54x54 matrix of pairwise link qualities
- Do PCA
- Project down to 2 principal dimensions
- PCA discovered the map of the lab



Problems and limitations of PCA

- What if very large dimensional data?
 - e.g., Images ($d \geq 10^4$)
- Problem:
 - Covariance matrix Σ is size (d^2)
 - $d=10_4 \rightarrow |\Sigma| = 10^8$
- Singular Value Decomposition (SVD)!
 - efficient algorithms available
 - some implementations find just top N eigenvectors


Laplacian Eigenmaps

Manifold


- A manifold is a **topological space** which is **locally Euclidean**. In general, any object which is nearly "flat" on small scales is a manifold.

Examples of 1-D manifolds include a line, a circle, and two separate circles.

sphere





torus



double torus




cross surface Klein Bottle



Embedding

An embedding is a **representation of a topological object**, manifold, graph, field, etc. in a certain space in such a way that its connectivity or algebraic properties are preserved.


Examples:



Embedding

An embedding is a **representation of a topological object**, manifold, graph, field, etc. in a certain space in such a way that its connectivity or algebraic properties are preserved.

Examples:



Embedding

An embedding is a **representation of a topological object**, manifold, graph, field, etc. in a certain space in such a way that its connectivity or algebraic properties are preserved.

Examples:



Integer



Real

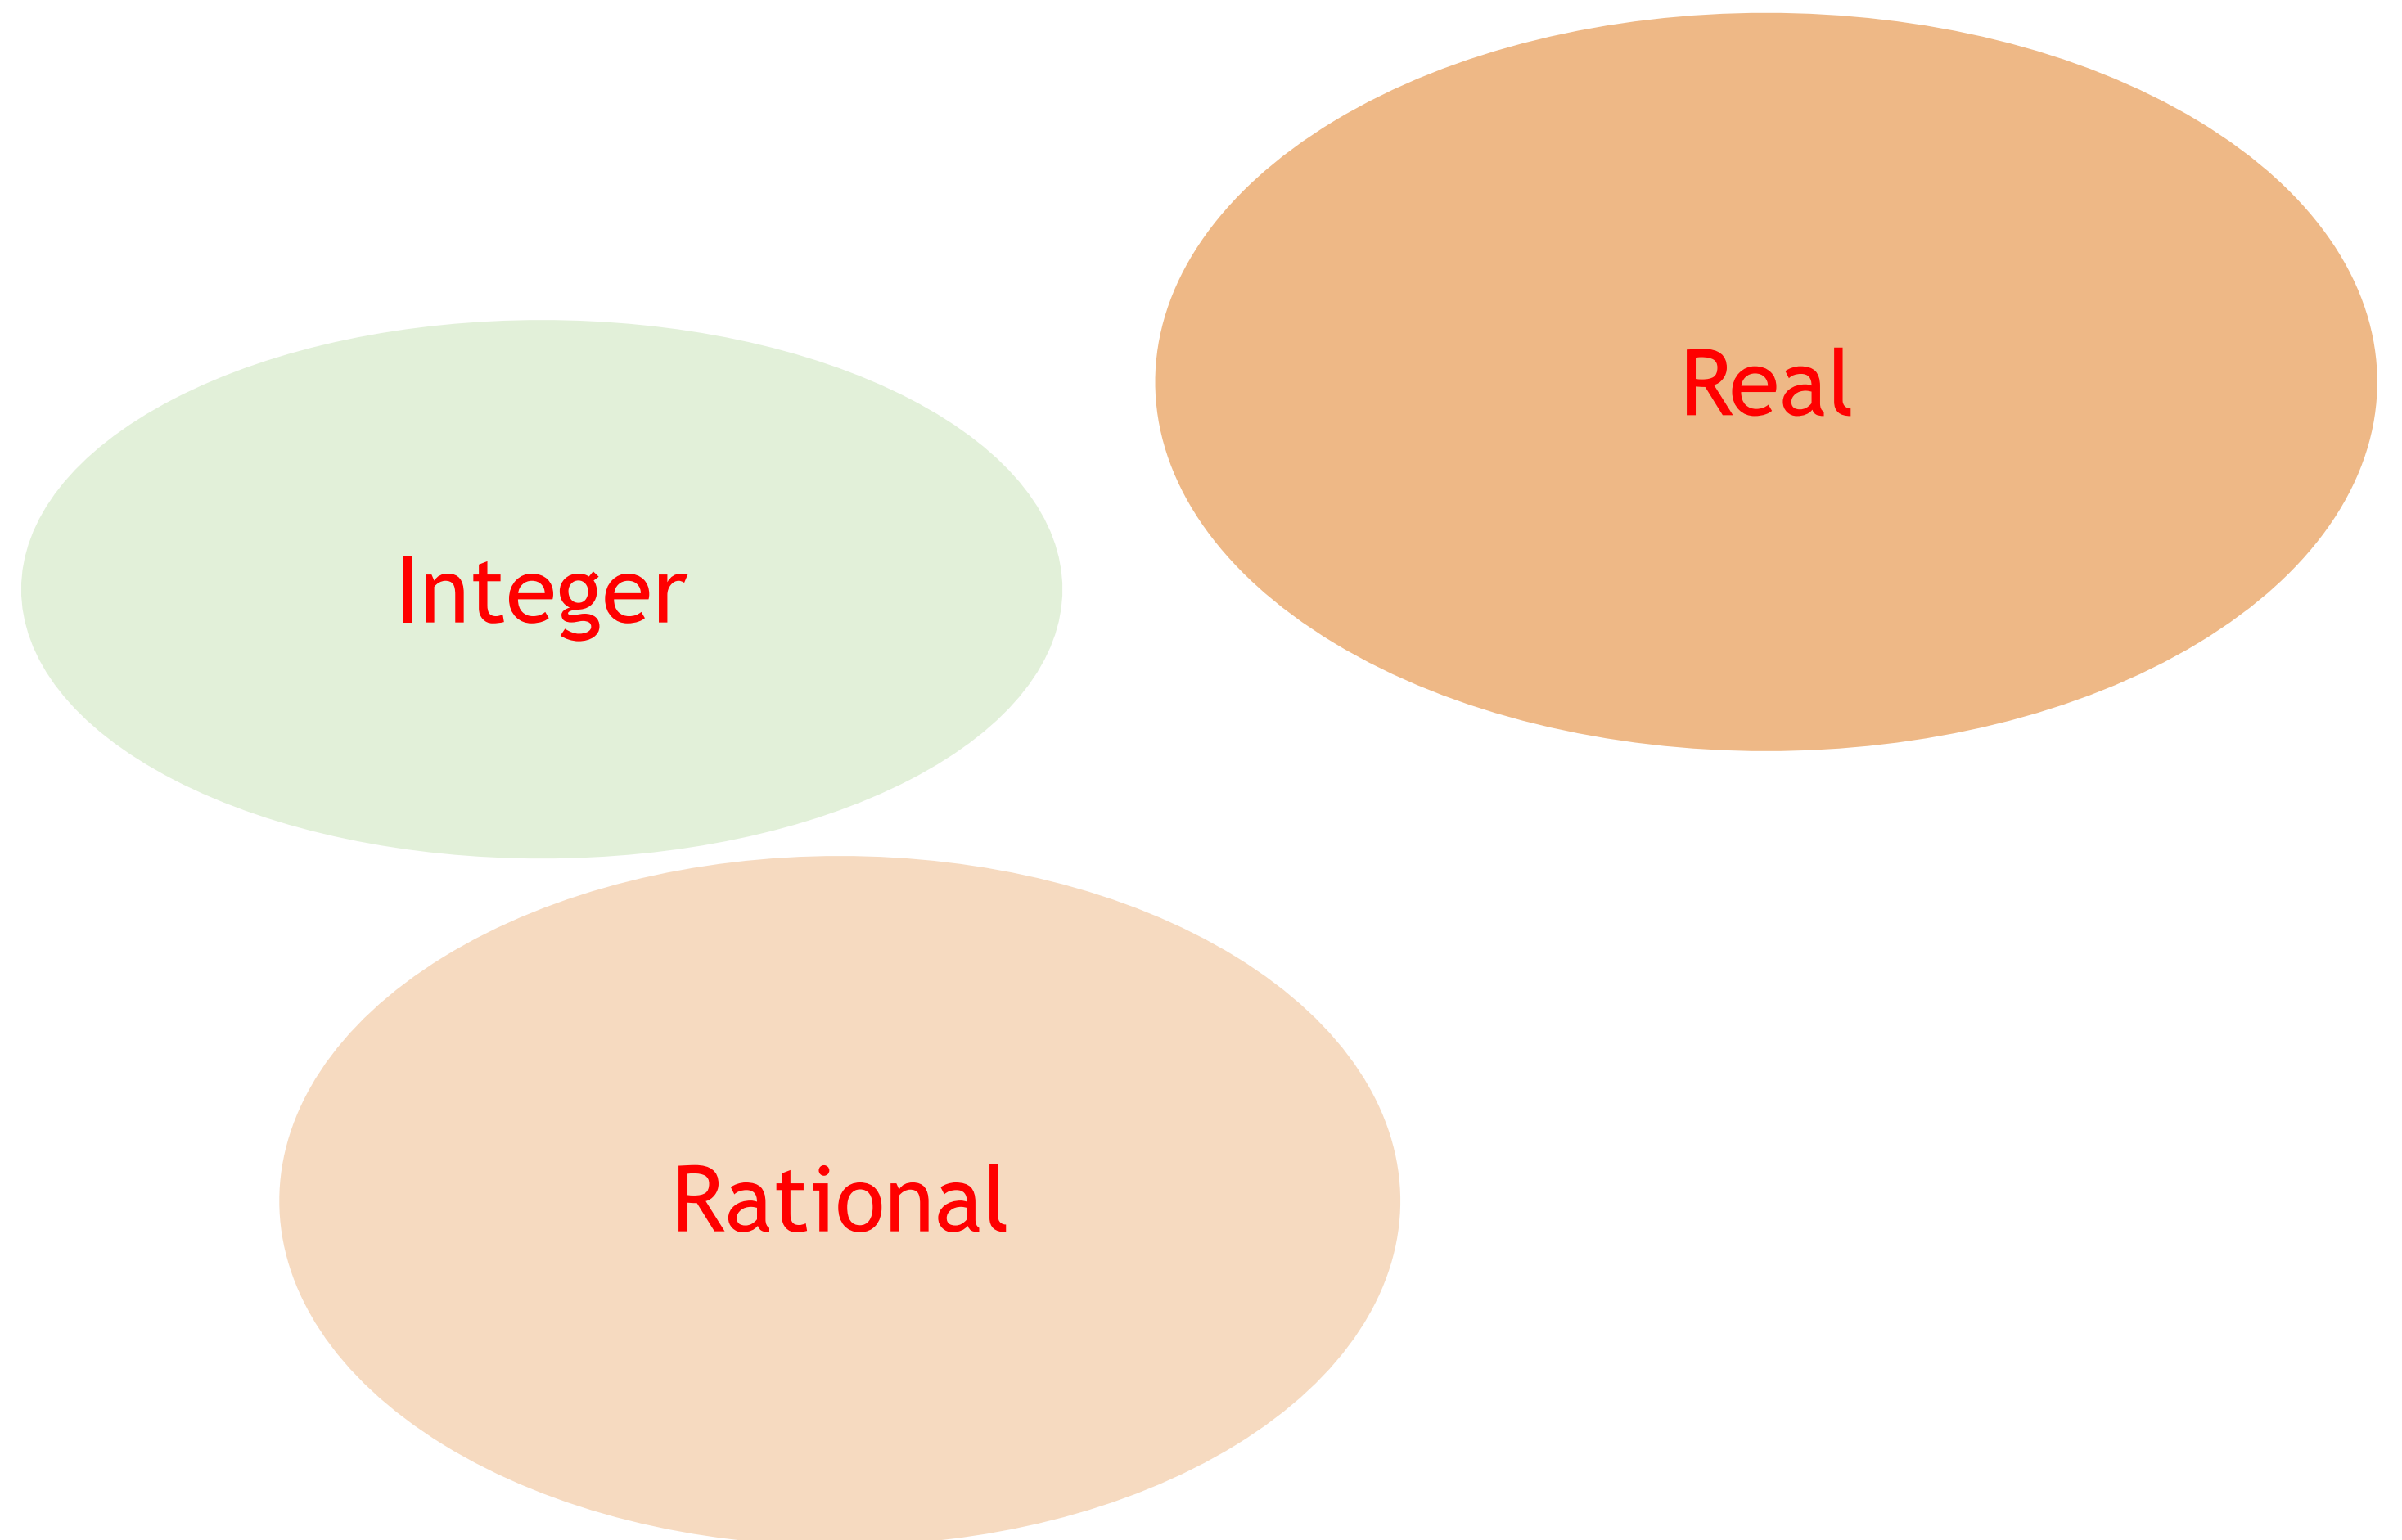


Rational

Embedding

An embedding is a **representation of a topological object**, manifold, graph, field, etc. in a certain space in such a way that its connectivity or algebraic properties are preserved.


Examples:



Embedding

An embedding is a **representation of a topological object**, manifold, graph, field, etc. in a certain space in such a way that its connectivity or algebraic properties are preserved.


Examples:



Embedding


An embedding is a **representation of a topological object**, manifold, graph, field, etc. in a certain space in such a way that its connectivity or algebraic properties are preserved.

Examples:




Manifold and Dimensionality Reduction

- Manifold: generalized “subspace” in \mathbf{R}^n
- Points in a local region on a manifold can be indexed by a subset of \mathbf{R}^k ($k \ll n$)




Manifold and Dimensionality Reduction

- Manifold: generalized “subspace” in \mathbf{R}^n
- Points in a local region on a manifold can be indexed by a subset of \mathbf{R}^k ($k \ll n$)




Manifold and Dimensionality Reduction

- Manifold: generalized “subspace” in \mathbf{R}^n
- Points in a local region on a manifold can be indexed by a subset of \mathbf{R}^k ($k \ll n$)




Manifold and Dimensionality Reduction

- Manifold: generalized “subspace” in \mathbf{R}^n
- Points in a local region on a manifold can be indexed by a subset of \mathbf{R}^k ($k \ll n$)




Manifold and Dimensionality Reduction

- Manifold: generalized “subspace” in \mathbf{R}^n
- Points in a local region on a manifold can be indexed by a subset of \mathbf{R}^k ($k \ll n$)




Manifold and Dimensionality Reduction

- Manifold: generalized “subspace” in \mathbf{R}^n
- Points in a local region on a manifold can be indexed by a subset of \mathbf{R}^k ($k \ll n$)



Manifold and Dimensionality Reduction


- If there is a global indexing scheme for M that maps a data point y on M



x : coordinate of z

Manifold and Dimensionality Reduction


- If there is a global indexing scheme for M that maps a data point y on M



x : coordinate of z

Manifold and Dimensionality Reduction


- If there is a global indexing scheme for M that maps a data point y on M



x : coordinate of z

Manifold and Dimensionality Reduction


- If there is a global indexing scheme for M that maps a data point y on M



x : coordinate of z

Manifold and Dimensionality Reduction


- If there is a global indexing scheme for M that maps a data point y on M



x : coordinate of z \rightarrow reduced dimension
representation of y


Geodesic Distance

- **Geodesic**: the shortest curve on a manifold that connects two points on the manifold
 - Example: on a sphere, geodesics are great circles
- **Geodesic distance**: length of the geodesic




Geodesic Distance

- **Geodesic**: the shortest curve on a manifold that connects two points on the manifold
 - Example: on a sphere, geodesics are great circles
- **Geodesic distance**: length of the geodesic




Geodesic Distance

- **Geodesic**: the shortest curve on a manifold that connects two points on the manifold
 - Example: on a sphere, geodesics are great circles
- **Geodesic distance**: length of the geodesic



Geodesic Distance


- **Geodesic**: the shortest curve on a manifold that connects two points on the manifold
 - Example: on a sphere, geodesics are great circles
- **Geodesic distance**: length of the geodesic



Geodesic Distance

- Euclidean distance **may not be** a good measure between two points on a manifold

- **Length of geodesic** is more appropriate




LLE and Laplacian Eigenmap

- The graph-based algorithms have **3 basic steps**:
 1. Find K nearest neighbors.
 2. Estimate local properties of manifold by looking at neighborhoods found in Step 1.
 3. Find a global embedding that preserves the properties found in Step 2.

Laplacian of a Graph

- Let $G(V,E)$ be a undirected graph without graph loops.
The Un-normalized Laplacian of the graph is


$$L=D-A$$



$$L = \begin{bmatrix} 3 & -1 & -1 & -1 \\ -1 & 2 & -1 & 0 \\ -1 & -1 & 2 & 0 \\ -1 & 0 & 0 & 1 \end{bmatrix} = \underbrace{\begin{bmatrix} 3 & 0 & 0 & 0 \\ 0 & 2 & 0 & 0 \\ 0 & 0 & 2 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}}_D - \underbrace{\begin{bmatrix} 0 & 1 & 1 & 1 \\ 1 & 0 & 1 & 0 \\ 1 & 1 & 0 & 0 \\ 1 & 0 & 0 & 0 \end{bmatrix}}_A$$


Laplacian Eigenmap

- Consider that X is a set of points in M , and M is a manifold embedded in \mathbb{R}^n .
- Find $\underline{y}_1, \dots, \underline{y}_n$ in \mathbb{R}^m such that \underline{y}_i represents $\underline{x}_i (m \ll n)$



Laplacian Eigenmap

- Construct the adjacency graph to approximate the manifold



Laplacian Eigenmap

There are two variations for W (weight matrix)

- **simple-minded** (1 if connected, 0 o.w.)
- **heat kernel** (t is real)

$$A_{ij} = e^{-\frac{\|x_i - x_j\|^2}{t}}$$

Laplacian Eigenmap N-Dimensional case

- Now we consider the more general problem of embedding the graph into **m-dimensional** Euclidean space
- Let Y be such a $n \times m$ map

$$Y = \begin{bmatrix} y_{11} & y_{12} & \cdots & y_{1m} \\ y_{21} & y_{22} & \cdots & y_{2m} \\ \vdots & \vdots & \cdots & \vdots \\ y_{n1} & y_{n2} & \cdots & y_{nm} \end{bmatrix}$$

- N-dimensional **dirichlet** energy

$$\operatorname{argmin}_Y \operatorname{trace}(Y'LY)$$

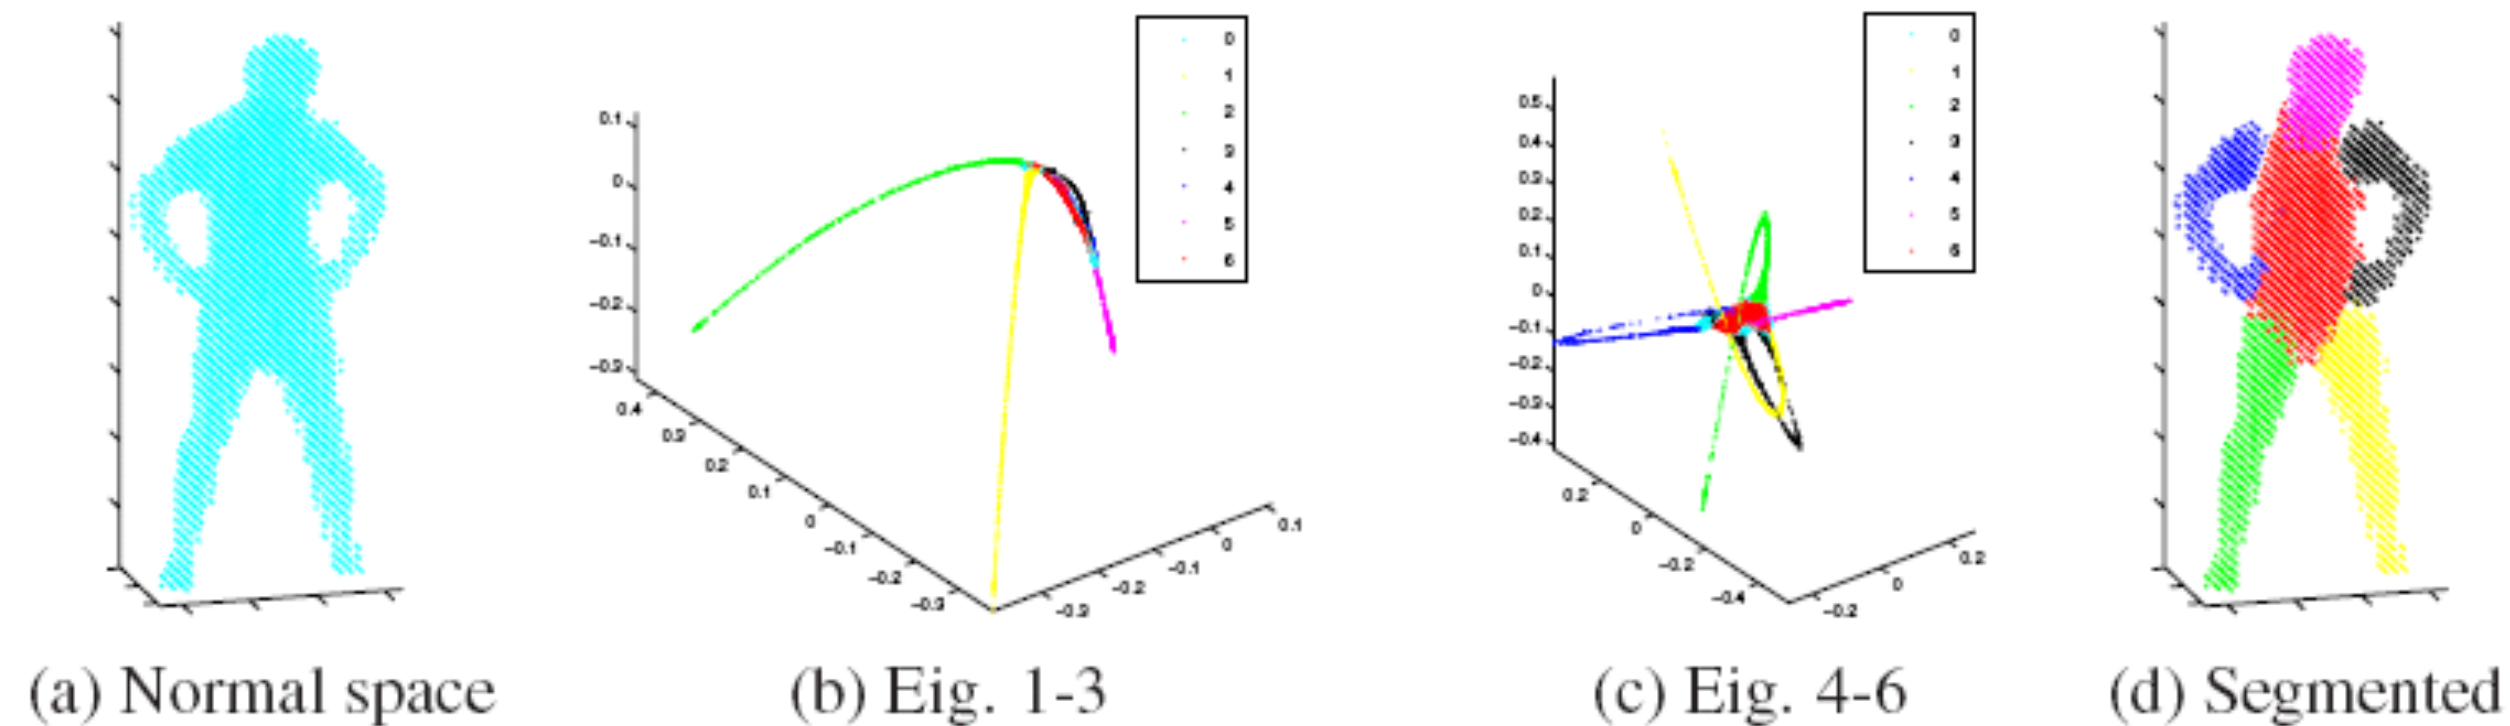
with

$$Y'Y = I$$


Solutions are the first m eigenvectors

Applications

- We can apply manifold learning to pattern recognition (face, handwriting etc)
- Recently, ISOMAP and Laplacian eigenmap are used to initialize the human body model.



Handwritten digit visualization



Considerations

- PROS:


Laplacian eigenmap provides a computationally efficient approach to non-linear dimensionality reduction that has locality preserving properties

BUT

Laplacian Eigenmap attempts to approximate or preserve neighborhood information


If you need GLOBAL consistency? -> Look at the ISOMAP method

T-SNE



T - SNE


- Map point from High Dimensional space (x)to low dimensional space (y) preserving points distributions
- density distribution around single points are preserved by the objective



$$p_{j|i} = \frac{\exp(-||x_i - x_j||^2 / 2\sigma_i^2)}{\sum_{k \neq i} \exp(-||x_i - x_k||^2 / 2\sigma_i^2)}$$

T - SNE

- Map point from High Dimensional space (x)to low dimensional space (y) preserving points distributions
- density distribution around single points are preserved by the objective



$$p_{j|i} = \frac{\exp(-||x_i - x_j||^2/2\sigma_i^2)}{\sum_{k \neq i} \exp(-||x_i - x_k||^2/2\sigma_i^2)}$$

Stochastic neighbour embedding

- Similarity of datapoint is converted to probabilities

$$p_{j|i} = \frac{\exp(-||x_i - x_j||^2 / 2\sigma_i^2)}{\sum_{k \neq i} \exp(-||x_i - x_k||^2 / 2\sigma_i^2)}$$

- Similarity in the low dimensional space y

$$q_{j|i} = \frac{\exp(-||y_i - y_j||^2)}{\sum_{k \neq i} \exp(-||y_i - y_k||^2)}$$

- OBJECTIVE:

Make the two distributions be as close as possible

- Minimize the Kullback Liebler Divergence

$$C = \sum_i KL(P_i || Q_i) = \sum_i \sum_j p_{j|i} \log \frac{p_{j|i}}{q_{j|i}}$$

Gradient descent solution

- Solve the problem pointwise by taking gradient of C w.r.t. points $i=1\dots n$

$$\frac{\partial C}{\partial y_i} = \sum_{j \neq i} (p_{j|i} - q_{j|i} + p_{i|j} - q_{i|j})(y_i - y_j)$$

1. Start with random y_i with $i=1\dots n$ with n number of points
2. Move y_i using gradient step update
3. DO it for all points in Y using momentum

$$Y^{(t)} = Y^{(t-1)} + \eta \frac{\partial C}{\partial y_i} + \beta(t)(Y^{(t-1)} - Y^{(t-2)})$$

Gradient proof here:



L.J.P. van der Maaten and G.E. Hinton. Visualizing High-Dimensional Data Using t-SNE. Journal of Machine Learning Research 9

https://lvdmaaten.github.io/publications/papers/JMLR_2008.pdf

Gradient Interpretation

- Similar to N body problem



$$\frac{\partial \mathcal{C}}{\partial y_i} = \sum_{j \neq i} (p_{j|i} - q_{j|i} + p_{i|j} - q_{i|j})(y_i - y_j)$$



Gradient Interpretation

- Similar to N body problem

$$\frac{\partial \mathcal{C}}{\partial y_i} = \sum_{j \neq i} (p_{j|i} - q_{j|i} + p_{i|j} - q_{i|j})(y_i - y_j)$$



Example Netflix movies

- More examples can be found here <https://lvdmaaten.github.io/tsne/>



T-sne code and additional resources

- T-SNE is the most popular embedding visualization method now.
- It is in most of the ML packages
- Inside SCIKIT LEARN
- Code and implementation for different languages here <https://lvdmaaten.github.io/tsne/>
- Sigma is crucial a good example on how sigma affect mapping <https://distill.pub/2016/misread-tsne/>
- Different TSNE variants : Symmetric, BH , Random Tree based