

Bayes

Machine Learning and Deep Learning
Lesson #4

Probability

- The world is a very uncertain place
- 30 years of Artificial Intelligence and Database research danced around this fact
- And then a few AI researchers decided to use some ideas from the eighteenth century

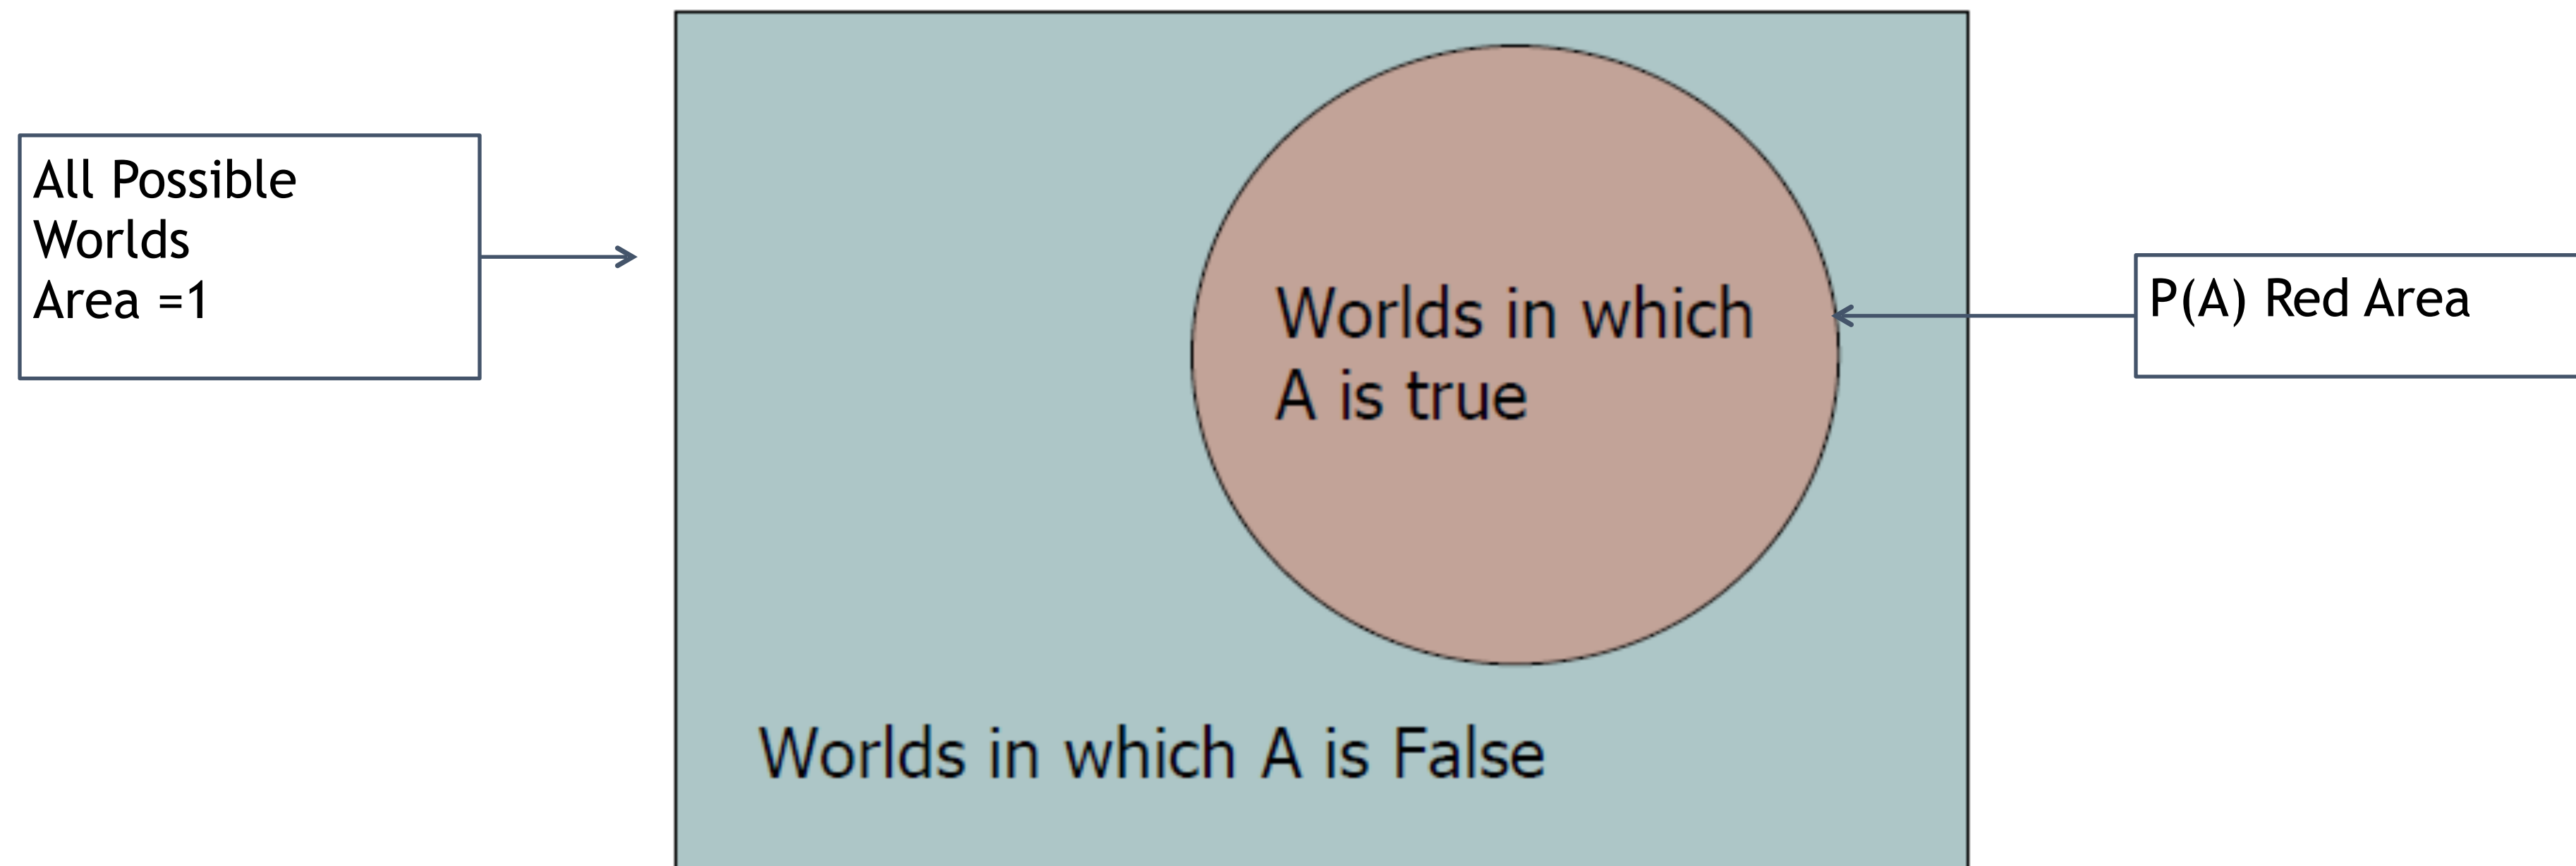


Discrete Random Variables

- A is a **Boolean-valued random variable** if A denotes an event, and there is some degree of uncertainty as to whether A occurs.
- Examples:
 - A = The US president in 2023 will be male
 - A = You wake up tomorrow with a headache
 - A = You have Ebola

Probabilities

- We write $P(A)$ as “the fraction of possible worlds in which A is true”
- We could at this point spend 2 hours on the philosophy of this. But we won't.



Multivalued Random Variables

- Suppose A can take on more than 2 values
- A is a **random variable** with *arity* k if it can take on exactly one value out of $\{v_1, v_2, \dots, v_k\}$

Thus:

$$P(A = v_i \wedge A = v_j) = 0 \text{ if } i \neq j$$
$$P(A = v_1 \vee A = v_2 \vee \dots \vee A = v_k) = 1$$

Properties:

- From the **Axioms of Probability** we can derive:

- Sum Rule:
$$P(A = v_1 \vee A = v_2 \vee A = v_i) = \sum_{j=1}^i P(A = v_j)$$

- Total Probability Rule:
$$\sum_{j=1}^k P(A = v_j) = 1$$

- Thus:
$$P(B \wedge [A = v_1 \vee A = v_2 \vee A = v_i]) = \sum_{j=1}^i P(B \wedge A = v_j)$$

- Discrete Marginalization over A:
$$P(B) = \sum_{j=1}^{\kappa} P(B \wedge A = v_j)$$

Conditional Probability

- Definition Conditional Probability:


$$P(A | B) = \frac{P(A, B)}{P(B)}$$

- Corollary Chain Rule:

$$P(A, B) = P(A | B) P(B)$$

Probabilistic Inference Problem

One day you wake up with a headache. You think: “Drat!
50% of flus are associated with headaches so I must have a
50-50 chance of coming down with flu”




$$\begin{aligned}P(H) &= 1/10 \\P(F) &= 1/40 \\P(H|F) &= 1/2\end{aligned}$$

Is this reasoning good?

Geometric Interpretation

Thanks to Jahanzeb Sherwani for contributing this explanation:



Let's say we have $P(F)$, $P(H)$, and $P(H|F)$, like in the example in class.

Areawise, $P(F) = A + B$, $P(H) = B + C$,

$$\text{Also, } P(H|F) = \frac{B}{A + B}$$

Thus, to get the opposite conditional probability, ie, $P(F|H)$, we need to figure out $\frac{B}{B + C}$


Since we know $B / (A+B)$, we can get $B / (B+C)$ by multiplying by $(A+B)$ and dividing by $(B+C)$. But since we already calculated, $A+B = P(F)$, and $B+C = P(H)$, so we are actually multiplying by $P(F)$ and dividing by $P(H)$. Which is Bayes Rule:

$$P(F|H) = P(H|F) * \frac{P(F)}{P(H)}$$

The Bayes Rule

- What we did geometrically?
- The **Bayes Rule**

$$P(B|A) = \frac{P(A, B)}{P(B)} = \frac{P(A|B) P(B)}{P(A)}$$




Bayes, Thomas (1763) An essay towards solving a problem in the doctrine of chances. *Philosophical Transactions of the Royal Society of London*, 53:370-418

«The intuition of a reverend of XVIII century changed the modern world and yours!!»



Joint Probability

- Two multivalued Random Variables A and B



- Inference**
“get insight about the occurrence of an Event from the JOINT”

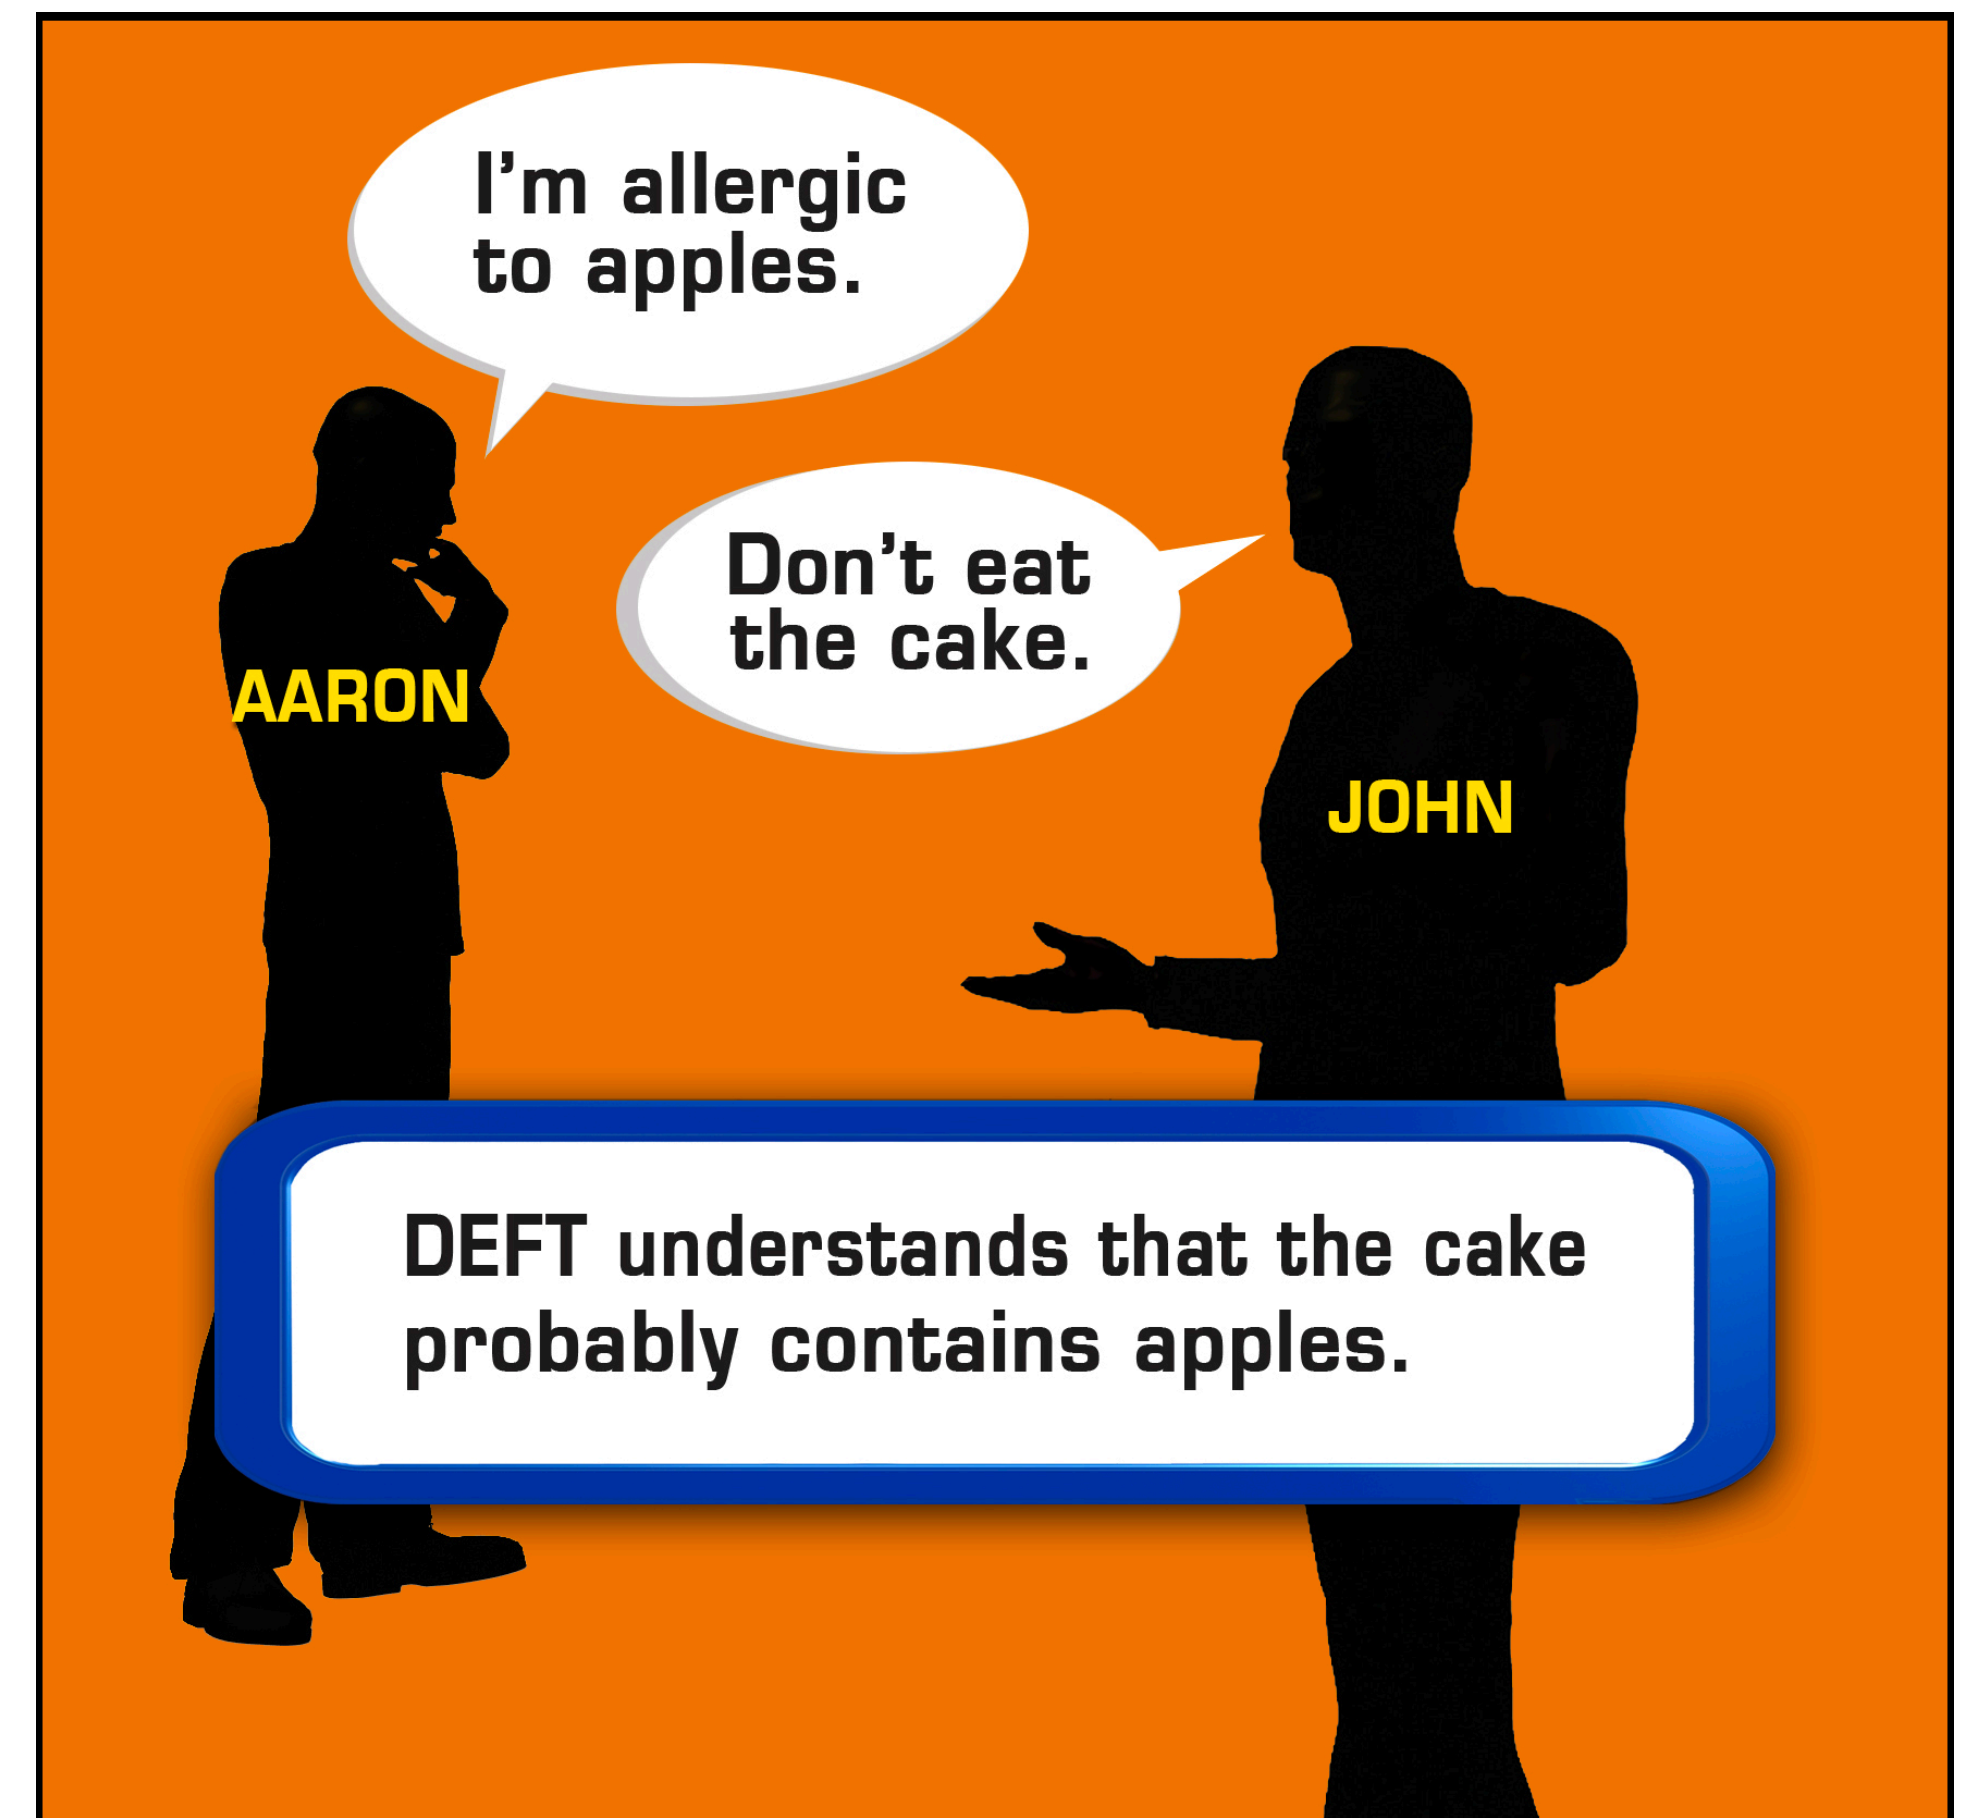
- E.g. if I work <40 what is the probability I am poor

$$P(\text{poor} | < 40) = \frac{P(\text{poor}, < 40)}{P(< 40)} = \frac{\sum_{\text{Male, Female}} P(\text{sex}, \text{poor}, < 40)}{P(< 40)} = \text{approx } 80 \%$$

Inference is a big deal

- I've got this evidence.
"What's the chance that this conclusion is true?"
- I've got a sore neck: how likely am I to have meningitis?
- There's a thriving set of industries growing based around Bayesian Inference.

Highlights are:
Medicine, Pharma, Help Desk Support, Engine Fault Diagnosis



How to compute Joint Probability

- **Idea One:** Expert Humans
- **Idea Two:** Simpler probabilistic facts and some algebra

Example: Suppose you knew

$$\begin{array}{ll} P(A) = 0.7 & P(C|A \wedge B) = 0.1 \\ & P(C|A \wedge \sim B) = 0.8 \\ P(B|A) = 0.2 & P(C|\sim A \wedge B) = 0.3 \\ P(B|\sim A) = 0.1 & P(C|\sim A \wedge \sim B) = 0.1 \end{array}$$

Then you can automatically
compute the JD using the
chain rule

$$\begin{array}{l} P(A=x \wedge B=y \wedge C=z) = \\ P(C=z|A=x \wedge B=y) P(B=y|A=x) P(A=x) \end{array}$$

In another lecture:
Bayes Nets, a
systematic way to
do this.

- **Idea Three:** Learn from Data

How to compute Joint Probability

- **Idea Two:** Simpler probabilistic facts and some algebra

Example: Suppose you knew

$$P(A) = 0.7 \quad P(C|A \wedge B) = 0.1$$

$$P(C|A \wedge \sim B) = 0.8$$

$$P(B|A) = 0.2 \quad P(C|\sim A \wedge B) = 0.3$$

$$P(B|\sim A) = 0.1 \quad P(C|\sim A \wedge \sim B) = 0.1$$

Then you can automatically
compute the JD using the
chain rule

$$P(A=x \wedge B=y \wedge C=z) = \\ P(C=z|A=x \wedge B=y) P(B=y|A=x) P(A=x)$$

In another lecture:
Bayes Nets, a
systematic way to
do this.

- **Idea Three:** Learn from Data

How to compute Joint Probability

- **Idea Three:** Learn from Data

Build a JD table for your attributes in which the probabilities are unspecified


A	B	C	Prob
0	0	0	?
0	0	1	?
0	1	0	?
0	1	1	?
1	0	0	?
1	0	1	?
1	1	0	?
1	1	1	?

The fill in each row with

$$\hat{P}(\text{row}) = \frac{\text{records matching row}}{\text{total number of records}}$$

A	B	C	Prob
0	0	0	0.30
0	0	1	0.05
0	1	0	0.10
0	1	1	0.05
1	0	0	0.05
1	0	1	0.10
1	1	0	0.25
1	1	1	0.10

Fraction of all records in which
A and B are True but C is False



How to compute Joint Probability

- **Idea Three:** Learn from Data

Build a JD table for your attributes in which the probabilities are unspecified


A	B	C	Prob
0	0	0	?
0	0	1	?
0	1	0	?
0	1	1	?
1	0	0	?
1	0	1	?
1	1	0	?
1	1	1	?

The fill in each row with

$$\hat{P}(\text{row}) = \frac{\text{records matching row}}{\text{total number of records}}$$

A	B	C	Prob
0	0	0	0.30
0	0	1	0.05
0	1	0	0.10
0	1	1	0.05
1	0	0	0.05
1	0	1	0.10
1	1	0	0.25
1	1	1	0.10

Fraction of all records in which
A and B are True but C is False



How to compute Joint Probability

- **Idea Three:** Learn from Data

Build a JD table for your attributes in which the probabilities are unspecified


A	B	C	Prob
0	0	0	?
0	0	1	?
0	1	0	?
0	1	1	?
1	0	0	?
1	0	1	?
1	1	0	?
1	1	1	?

The fill in each row with

$$\hat{P}(\text{row}) = \frac{\text{records matching row}}{\text{total number of records}}$$

A	B	C	Prob
0	0	0	0.30
0	0	1	0.05
0	1	0	0.10
0	1	1	0.05
1	0	0	0.05
1	0	1	0.10
1	1	0	0.25
1	1	1	0.10

Fraction of all records in which
A and B are True but C is False



How to compute Joint Probability

- **Idea Three:** Learn from Data

Build a JD table for your attributes in which the probabilities are unspecified


A	B	C	Prob
0	0	0	?
0	0	1	?
0	1	0	?
0	1	1	?
1	0	0	?
1	0	1	?
1	1	0	?
1	1	1	?

The fill in each row with

$$\hat{P}(\text{row}) = \frac{\text{records matching row}}{\text{total number of records}}$$

A	B	C	Prob
0	0	0	0.30
0	0	1	0.05
0	1	0	0.10
0	1	1	0.05
1	0	0	0.05
1	0	1	0.10
1	1	0	0.25
1	1	1	0.10

Fraction of all records in which
A and B are True but C is False



How to compute Joint Probability

- **Idea Three:** Learn from Data

Build a JD table for your attributes in which the probabilities are unspecified


A	B	C	Prob
0	0	0	?
0	0	1	?
0	1	0	?
0	1	1	?
1	0	0	?
1	0	1	?
1	1	0	?
1	1	1	?

The fill in each row with

$$\hat{P}(\text{row}) = \frac{\text{records matching row}}{\text{total number of records}}$$

A	B	C	Prob
0	0	0	0.30
0	0	1	0.05
0	1	0	0.10
0	1	1	0.05
1	0	0	0.05
1	0	1	0.10
1	1	0	0.25
1	1	1	0.10

Fraction of all records in which
A and B are True but C is False



How to compute Joint Probability

Build a JD table for your attributes in which the probabilities are unspecified


A	B	C	Prob
0	0	0	?
0	0	1	?
0	1	0	?
0	1	1	?
1	0	0	?
1	0	1	?
1	1	0	?
1	1	1	?

The fill in each row with

$$\hat{P}(\text{row}) = \frac{\text{records matching row}}{\text{total number of records}}$$


A	B	C	Prob
0	0	0	0.30
0	0	1	0.05
0	1	0	0.10
0	1	1	0.05
1	0	0	0.05
1	0	1	0.10
1	1	0	0.25
1	1	1	0.10

Fraction of all records in which
A and B are True but C is False



Density Estimation

- Our Joint Distribution learner is our first example of something called **Density Estimation**
- A **Density Estimator** learns a mapping from a set of attributes to a Probability



- Density estimation can be:
 - **Observing variables values**: Discrete/Continuous
 - **Observing probability equation**: Parametric/Non Parametric

Density Estimation Evaluation

- Given a record \mathbf{x} , a density estimator M can tell you how likely the record is

$$\hat{P}(\mathbf{x}|M)$$

- Given a dataset with R records the DE can tell you how likely the dataset is
 - (assuming data independently generated from DE JD)

$$\hat{P}(\text{dataset}|M) = \hat{P}(\mathbf{x}_1 \wedge \mathbf{x}_2 \dots \wedge \mathbf{x}_R|M) = \prod_{k=1}^R \hat{P}(\mathbf{x}_k|M)$$

- Since probabilities of datasets get so small we usually use **log** probabilities

$$\log \hat{P}(\text{dataset}|M) = \log \prod_{k=1}^R \hat{P}(\mathbf{x}_k|M) = \sum_{k=1}^R \log \hat{P}(\mathbf{x}_k|M)$$

Density Estimators Pros

- We have a way to learn a Density Estimator from data.
- Density estimators can do many good Things:
 - Can sort the records by probability, and thus spot weird records (anomaly detection)
 - Can do inference: $P(E1|E2)$ (Automatic Doctor / Help Desk etc)

Density Estimators Pros

- We have a way to learn a Density Estimator from data.
- Density estimators can do many good Things:
 - Can sort the records by probability, and thus spot weird records (anomaly detection)
 - Can do inference: $P(E1|E2)$ (Automatic Doctor / Help Desk etc)

BUT

**Density estimation by directly
learning the joint is trivial, mindless
and dangerous**

Overfitting

If **this** ever happens, it means there are certain combinations that we learn are impossible

mpg	modelyear	maker		
bad	70to74	america	0.27551	<div></div>
		asia	0.0255102	<div></div>
		europa	0.0153061	<div></div>
	75to77	america	0.153061	<div></div>
		asia	0.0255102	<div></div>
		europa	0.0357143	<div></div>
	78to83	america	0.0561224	<div></div>
		asia	Never	
		europa	Never	
good	70to74	america	0.0102041	<div></div>

$$\begin{aligned}\log \hat{P}(\text{testset}|M) &= \log \prod_{k=1}^R \hat{P}(\mathbf{x}_k|M) = \sum_{k=1}^R \log \hat{P}(\mathbf{x}_k|M) \\ &= -\infty \text{ if for any } k \hat{P}(\mathbf{x}_k|M) = 0\end{aligned}$$

Overfitting

If **this** ever happens, it means there are certain combinations that we learn are impossible

mpg	modelyear	maker		
bad	70to74	america	0.27551	<div></div>
		asia	0.0255102	<div></div>
		europa	0.0153061	<div></div>
	75to77	america	0.153061	<div></div>
		asia	0.0255102	<div></div>
		europa	0.0357143	<div></div>
	78to83	america	0.0561224	<div></div>
		asia	Never	
		europa	Never	
good	70to74	america	0.0102041	<div></div>

$$\begin{aligned}\log \hat{P}(\text{testset}|M) &= \log \prod_{k=1}^R \hat{P}(\mathbf{x}_k|M) = \sum_{k=1}^R \log \hat{P}(\mathbf{x}_k|M) \\ &= -\infty \text{ if for any } k \hat{P}(\mathbf{x}_k|M) = 0\end{aligned}$$

We need Density Estimators that are less prone to overfitting

Overfitting



Naive Density Estimator

- The problem with the Joint Estimator is that it just mirrors the training data.
- We need something which **generalizes** more usefully.

The **naïve model** generalizes strongly:

“Assume that each attribute is distributed independently of any of the other attributes.”

IID Independently Distributed Data

- Let $x[i]$ denote the i -th field of record x .
- The independently distributed assumption says that:

for any $i, v, u_1, u_2 \dots u_{i-1}, u_{i+1} \dots U_m$

- $x[i]$ is **independent** of $\{x[1], x[2], \dots x[i-1], x[i+1], \dots x[M]\}$

$$x[i] \perp \{x[1], x[2], \dots x[i-1], x[i+1], \dots x[M]\}$$

Independence Theorems

- Given A and B random variables
- A is **independent** of B «if and only if» $P(A|B)=P(A)$

Consequences:

- $P(A,B)=P(A)P(B)$
- $P(B|A)=P(B)$
- $P(\sim A|B)=P(\sim A)$
- $P(A|\sim B)=P(A)$

Naive DE General Case


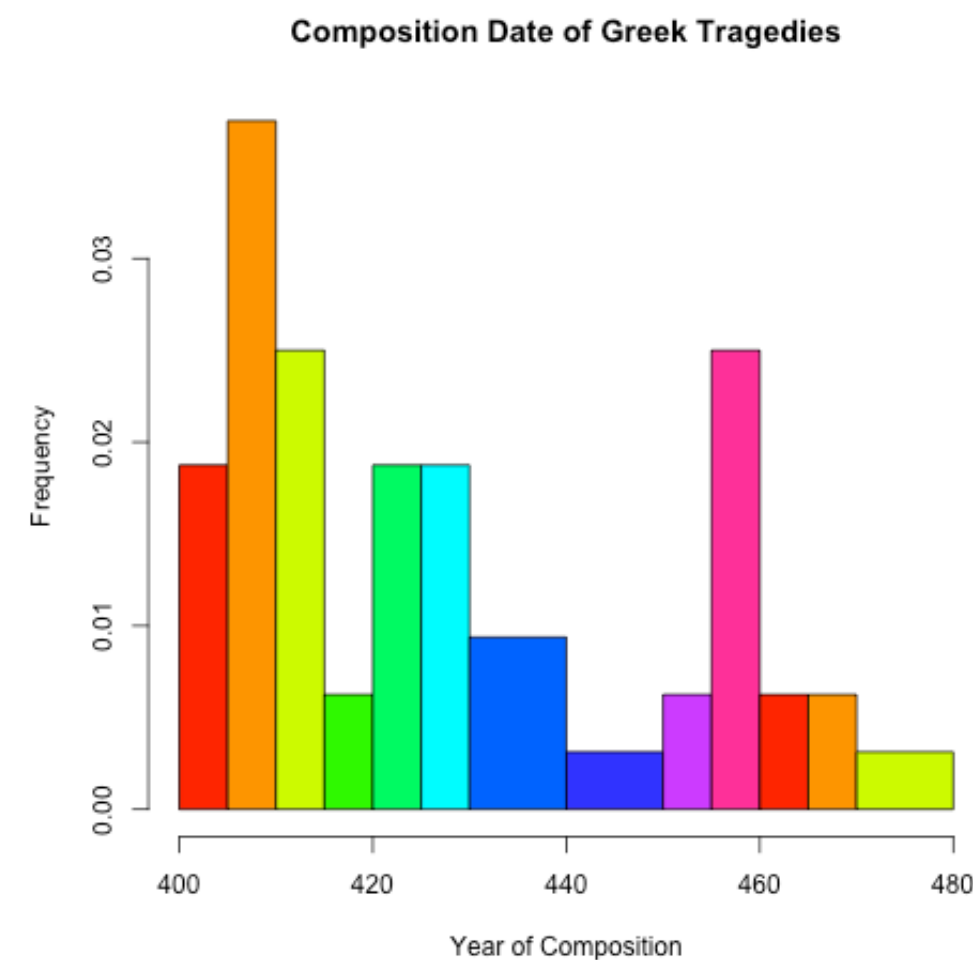

- Suppose $x[1], x[2], \dots, x[M]$ are independently distributed.

$$P(x[1] = u_1, x[2] = u_2, \dots, x[M] = u_M) = \prod_{k=1}^M P(x[k] = u_k)$$

But How do we learn a naïve density estimator:

$$\hat{P}(x[i] = u) = \frac{\text{\#records in which } x[i] = u}{\text{total number of records}}$$

- **Normalized Histogram** is a discrete Non Parametric DE



Bayes Classifier

Build a Bayes Classifier (Preliminary Step)

1. Assume you want to predict output Y which has arity n_Y and values V_1, V_2, \dots, V_{n_Y}
3. Assume there are m input attributes called X_1, X_2, \dots, X_m
5. Break dataset into n_Y smaller datasets called $DS_1, DS_2, \dots, DS_{n_Y}$
7. Define $DS_i = \text{Records in which } Y=v_i$
9. For each DS_i learn Density Estimator M_i to model the input distribution among the $Y=v_i$ records.

Build a Bayes Classifier (Preliminary Step)

1. Assume you want to predict output Y which has arity n_Y and values V_1, V_2, \dots, V_{n_Y}
3. Assume there are m input attributes called X_1, X_2, \dots, X_m
5. Break dataset into n_Y smaller datasets called $DS_1, DS_2, \dots, DS_{n_Y}$
7. Define $DS_i = \text{Records in which } Y=v_i$
9. For each DS_i learn Density Estimator M_i to model the input distribution among the $Y=v_i$ records.

M_i estimates

$$P(X_1, X_2, \dots, X_m \mid Y=v_i)$$

ML Classifier

- Idea: When a new set of input values ($X_1 = u_1, X_2 = u_2, \dots, X_m = u_m$) come along to be evaluated predict the value of Y that makes $P(X_1, X_2, \dots, X_m \mid Y = v_i)$ most likely

$$Y^{\text{predict}} = \underset{v}{\operatorname{argmax}} P(X_1 = u_1 \cdots X_m = u_m \mid Y = v)$$

ML Classifier

- Idea: When a new set of input values ($X_1 = u_1, X_2 = u_2, \dots, X_m = u_m$) come along to be evaluated predict the value of Y that makes $P(X_1, X_2, \dots, X_m \mid Y = v_i)$ most likely

$$Y^{\text{predict}} = \underset{v}{\operatorname{argmax}} P(X_1 = u_1 \cdots X_m = u_m \mid Y = v)$$

Is this a good idea?

ML Classifier

- Idea: When a new set of input values ($X_1 = u_1, X_2 = u_2, \dots, X_m = u_m$) come along to be evaluated predict the value of Y that makes $P(X_1, X_2, \dots, X_m \mid Y = v_i)$ most likely

$$Y^{\text{predict}} = \underset{v}{\operatorname{argmax}} P(X_1 = u_1 \cdots X_m = u_m \mid Y = v)$$

Is this a good idea?

$$Y^{\text{predict}} = \underset{v}{\operatorname{argmax}} P(X_1 = u_1 \cdots X_m = u_m \mid Y = v)$$

This is a **Maximum Likelihood Classifier**

Cons:

- Not Bayesian
- Silly if some Y_i are unlikely

Build a Bayes Classifier

- Much Better Idea!!!:
- When a new set of input values ($X_1 = u_1, X_2 = u_2, \dots, X_m = u_m$) come along to be evaluated predict the value of Y that makes most likely

$$P(Y=v_i | X_1, X_2, \dots, X_m)$$

$$Y^{\text{predict}} = \underset{v}{\operatorname{argmax}} P(Y = v | X_1 = u_1 \cdots X_m = u_m)$$

We can get the posterior using Bayes Rule

$$\begin{aligned} & P(Y = v | X_1 = u_1 \cdots X_m = u_m) \\ &= \frac{P(X_1 = u_1 \cdots X_m = u_m | Y = v)P(Y = v)}{P(X_1 = u_1 \cdots X_m = u_m)} \\ &= \frac{P(X_1 = u_1 \cdots X_m = u_m | Y = v)P(Y = v)}{\sum_{j=1}^{n_Y} P(X_1 = u_1 \cdots X_m = u_m | Y = v_j)P(Y = v_j)} \end{aligned}$$

Naive Version Bayes Classifiers

- Hypothize X are **independent** and use product rule to build the joint DE

$$Y^{\text{predict}} = \underset{v}{\operatorname{argmax}} P(Y = v) \prod_{j=1}^{n_Y} P(X_j = u_j \mid Y = v)$$

- Technical Hint: If you have 10,000 input attributes that product will underflow in floating point math. You should use **logs**.

Naive Version Bayes Classifiers

- Hypothize X are **independent** and use product rule to build the joint DE

$$Y^{\text{predict}} = \underset{v}{\operatorname{argmax}} P(Y = v) \prod_{j=1}^{n_Y} P(X_j = u_j | Y = v)$$

- Technical Hint: If you have 10,000 input attributes that product will underflow in floating point math. You should use **logs**.


$$Y^{\text{predict}} = \underset{v}{\operatorname{argmax}} \left(\log P(Y = v) + \sum_{j=1}^{n_Y} \log P(X_j = u_j | Y = v) \right)$$

Example Digit Recognition

Naïve Bayes Training

- Now that we've decided to use a Naïve Bayes classifier, we need to train it with some data:

MNIST Training Data



Naïve Bayes Training

- Training in Naïve Bayes is easy:
 - Estimate $P(Y=v)$ as the fraction of records with $Y=v$
 - Estimate $P(X_i=u|Y=v)$ as the fraction of records with $Y=v$ for which $X_i=u$

$$P(Y = v) = \frac{\text{Count}(Y = v)}{\# \text{ records}}$$

- (This corresponds to Maximum Likelihood estimation of model parameters)

$$P(X_i = u|Y = v) = \frac{\text{Count}(X_i = u \wedge Y = v)}{\text{Count}(Y = v)}$$

Naïve Bayes Training



- In practice, some of these counts can be zero
- Fix this by adding “virtual” counts:

$$P(X_i = u|Y = v) = \frac{\text{Count}(X_i = u \wedge Y = v) + 1}{\text{Count}(Y = v) + 2}$$




- (This is like putting a prior on parameters and doing MAP estimation instead of MLE)
- This is called *Smoothing*

Naïve Bayes Training

- For binary digits, training amounts to averaging all of the training fives together and all of the training sixes together.



Naïve Bayes Classification



Performance on a Test Set

- Naïve Bayes is often a good choice if you don't have much training data!

